

Sentiment analysis

I. Introduction

This project approaches an active research area in Natural Language Processing (NLP), sentiment analysis. Given the exponentially growing of online review data (Amazon, IMDB, etc...), sentiment analysis becomes increasingly important. We are going to build a sentiment classifier, i.e., evaluating a piece of text being either positive or negative.

This assignment requires to gather the raw data, do preprocessing, design suitable ML algorithms and implement the solution.

The "Large Movie Review Dataset"(1) is used for this project.

The dataset is compiled from a collection of 50,000 reviews from IMDB on the condition there are no more than 30 reviews per movie.

Number of positive and negative reviews are equal.

Negative reviews have scores lesser or equal 4 out of 10 while a positive review greater or equal 7 out of 10. Neutral reviews are not included.

Then, 50,000 reviews are divided evenly into the training and test sets

We are then going to train a Stochastic Gradient Descent (SGD) Classifier. While gradient descend is powerful, it can be prohibitively expensive when the dataset is extremely large because every single data point needs to be processed.

However, it turns out that, when the data is large, rather than the entire dataset, SGD algorithm performs just as good with a small random subset of the original data. This is the central idea of Stochastic SGD and particularly handy for the text data since corpus are often humongous.

(1) <https://ai.stanford.edu/amaas/data/sentiment/>

II. Methodology

1. Data Preprocessing

First the raw database used to train our classifier is combined into a single .csv file which has three columns, *row_number*, *text* and *polarity*. The column *text* contains review texts from the aclImdb database while the column *polarity* consists of sentiment labels, *1* for positive and *0* for negative.

Similarly the test database used to test our prediction model is a single csv file with two columns: *row_number* and *text*. The column *polarity* is excluded and the objective is to use the trained SGD classifier to predict this information.

In addition, common English stopwords are removed as well as any numbers and punctuation found.

Furthermore, words in the test database that do not appear in the training database are removed in order to avoid any mismatched features.

2. Data Representations

The very first step in solving any NLP problem is finding a way to represent the text data so that machines can understand. A common approach is to use a document-term vector where each document is encoded as a discrete vector that counts occurrences of each word in the vocabulary it contains. This data representation is also called a Unigram model.

For example, consider two one-sentence documents:

d1: "I love this Artificial Intelligence course."

d2: "Artificial Intelligence is awesome"

The vocabulary {artificial, awesome, this, course, I, intelligence, is, love} in the two documents can be encoded as v1 and v2 as follow:

	artificial	awesome	this	course	I	intelligence	is	love
v1	1	0	1	1	1	1	0	1
v2	1	1	0	0	0	1	1	0

Table 1: Unigram Data Representation

A more sophisticated data representation model is the Bigram model where occurrences depend on a sequence of two words rather than an individual one. Using the same example as before, v1 and v2 are now encoded as seen in Table 2:

	artificial artificial	artificial awesome	...	intelligence artificial	...	love love
v1	0	0	...	1	...	0
v2	0	1	...	0	...	0

Table 2: Bigram Data Representation

Sometimes, a very high word counting may not be meaningful. For example, a common word like “say” may appear 10 times more than a less-common word such as “machine” but it does not mean “say” is 10 times more relevant to our sentiment classifier. To alleviate this issue, we can instead use the term frequency $tf[t]$:

$$tf[t] = 1 + \log(f[t, d])$$

where $f[t, d]$ is the count of term t in document d . The log function dampens the unwanted influence of common English words.

Inverse document frequency (idf) is a similar concept. For example, Computer Science documents often have words such as computers, CPU, programming, appearing over and over. While they are not common English words, because of the document domain, their occurrences are very high. To rectify this, we can adjust using inverse term frequency $idf[t] = \log(N / df[t])$ where $df[t]$ is the number of documents containing the term t and N is the total number of document in the dataset.

Therefore, instead of just using a word frequency, a frequency tf_idf for each term t can be used:

$$tf_idf[t] = tf[t] * idf[t]$$

3. Training and testing of the classifier

Using the sklearn library in Python, the following steps are done for each data representation model:

- Step 1: Each text document is converted into one of the token-count row of a term-document matrix [number of samples x terms frequencies] .
- Step 2: A fitting of the SGD classifier is done by the minimization of a linear SVM objective function with a l1 penalty term (to bring sparsity to the model). Because SGD only sums on a random portion of the dataset, it is way less computationally extensive than GD even if more updates are necessary before convergence.
- Step 3: The SGD classifier can then be used to predict new polarity labels from test data.

III. Results

Four different data representation models were used and a SGD classifier was trained then used to predict sentiment polarity of new test data (whose polarities were in fact known). Table 3 presents a comparison between the prediction efficiency rates obtained for each representation. One observes that the best results are obtained with a model using a Unigram idf model, but that efficiencies between representations only differ by at most 2%. The advantage of using one type of representation model over another is thus not demonstrated here.

	Unigram	Bigram	Unigram idf	Bigram idf
prediction score	85.20%	85.75%	87.26%	86.50%

Table 3: Comparison of predictions between 4 different data representation models

IV. Python Code

```
1 from __future__ import division
2 import sys
3 from sklearn import linear_model
4 from sklearn import datasets
5 from sklearn.linear_model import SGDClassifier
6 from sklearn.model_selection import train_test_split
7 from sklearn.model_selection import cross_val_score
8 from sklearn.metrics import accuracy_score
9 from sklearn import datasets
10 from sklearn.linear_model import LogisticRegression
11 import numpy as np
12 import pandas as pd
13 from sklearn.feature_extraction.text import CountVectorizer
14 from sklearn.feature_extraction.text import TfidfTransformer
15 from sklearn.feature_extraction.text import TfidfVectorizer
16 import os
17 import os.path
18 import glob
19 import fileinput
20 import fnmatch
21 import time
22 import re
23 import gc
24
25 # train_path = "../resource/lib/publicdata/aclImdb/train/" # use terminal to ls
26 # test_path = "../resource/lib/publicdata/imdb_te.csv" # test data for grade eval
27
28 begin = time.clock()
29 train_path = "C:/MOOC/Edx-MOOC/AI micromaster/AI/Projects/Project5/aclImdb/train"
30 test_path = "C:/MOOC/Edx-MOOC/AI micromaster/AI/Projects/Project5/imdb_te.csv"
31
32 stopwords = set()
33 for line in open("stopwords.en.txt", 'r'):
34     stopwords.add(line[:-1])
35
36 stopwords2 = {'\'s', '\t', '\m', '\re', '\ ', '/', '\ll',
37              '1', '2', '3', '4', '5', '6', '7', '8', '9', '0', 'I', 'My', 'You', 'They',
38              'The', 'There', 'That', 'This ', 'Those', 'These',
39              'A', 'An', 'It', 'Or', ' ', ''}
40
41
42 ## Use nltk library for NLP
43 # from nltk.corpus import stopwords
44 # # ...
45 # filtered_words = [word for word in word_list if word not in stopwords.words('en')]
46
47
48
49
50 def imdb_data_preprocess(inpath, outpath=".", name="imdb_tr.csv", mix=False):
51     """Implement this module to extract
```

```

52 and combine text files under train_path directory into
53 imdb_tr.csv. Each text file in train_path should be stored
54 as a row in imdb_tr.csv. And imdb_tr.csv should have two
55 columns, "text" and label""
56
57 posi = os.listdir(inpath + '/pos')
58 nega = os.listdir(inpath + '/neg')
59
60 csv_file = pd.DataFrame(columns=['text', 'polarity'])
61
62 print(len(nega))
63
64 for i in range(len(posi)): # len(posi)
65     file = open(inpath + '/pos/' + posi[i])
66     text = file.read()
67     text1 = re.split('\W+', text)
68     text2 = text1[:]
69     for word in text1:
70         if word in stopwords:
71             text2.remove(word)
72         if word in stopwords2:
73             text2.remove(word)
74     text = ""
75     for word in text2:
76         text += word + ' '
77     csv_file.loc[i] = [text, 1]
78
79 for i in range(len(nega)): # len(nega)
80     file = open(inpath + '/neg/' + nega[i])
81     text = file.read()
82     text1 = re.split('\W+', text)
83     text2 = text1[:]
84     for word in text1:
85         if word in stopwords:
86             text2.remove(word)
87         if word in stopwords2:
88             text2.remove(word)
89     text = ""
90     for word in text2:
91         text += word + ' '
92     csv_file.loc[i + len(posi)] = [text, 0]
93
94 csv_file.to_csv(outpath + name, sep=',')
95
96 def prepare_data(train_data):
97     # prepare train data
98     fi_tr = train_data.iloc[:, 1]
99     pola = train_data.iloc[:, 2]
100     db0_tr = fi_tr.values
101     pola0_tr = pola.values
102     print('pola_tr:%s' % pola0_tr.shape)
103     print('db_tr:%s' % db0_tr.shape)
104     pola0_tr = pola0_tr.astype(int)

```

```

105     print('type_pola: %s'%type(pola0_tr[2]))
106     print('type_db_tr: %s' % type(db0_tr[2]))
107     # split train data for cross_validation
108     db_tr, X_test, pola_tr, y_test = train_test_split(db0_tr, pola0_tr, test_size=0.2)
109     print(len(db_tr), len(X_test))
110     # prepare test data
111     test_data = pd.read_csv("imdb_te.csv")
112     fi_te = test_data.iloc[:, 1]
113     db_te = fi_te.values
114     db_te2 = []
115     for line in db_te: # len(nega)
116         text1 = re.split('\W+', line)
117         text = ""
118         text2 = text1[:]
119         for word in text1:
120             if word in stopwords:
121                 text2.remove(word)
122             if word in stopwords2:
123                 text2.remove(word)
124
125         for word in text2:
126             text += word + ' '
127         db_te2.append(text)
128
129     return db_tr, db_te2, X_test, pola_tr, y_test
130
131 ## Unigram
132
133 def unigram(db_tr, db_te2, X_test, pola_tr, y_test):
134     vect = CountVectorizer()
135     vect.fit(db_tr)
136     Xtr = vect.transform(db_tr)
137     print('Xtr:(%s, %s)'%(Xtr.shape))
138     Xte = vect.transform(db_te2)
139     print('Xte:(%s, %s)'%(Xte.shape))
140     clf = SGDClassifier(loss='hinge', penalty='l1')
141     clf.fit(Xtr, pola_tr)
142     pred1 = clf.predict(Xte)
143     X_test1 = vect.transform(X_test)
144     print('score Unigram, test data: %s' % clf.score(Xte, pred1))
145     print('cross validation score Unigram, test data: %s' % cross_val_score(clf,
146     print('score Unigram: %s'%clf.score(X_test1, y_test))
147     print('cross validation score Unigram: %s'%cross_val_score(clf, X_test1, y_test))
148     # print(accuracy_score(y_test, pred1))
149
150     out = open('unigram.output.txt', 'w')
151     for i in range(len(pred1)):
152         out.write('%s\n'%int(pred1[i]))
153     out.close()
154
155     return Xtr, Xte, X_test1
156
157 ## bigram

```

```

158
159 def bigram(db_tr, db_te2, X_test, pola_tr, y_test):
160     bigram_vect = CountVectorizer(ngram_range=(1, 2), token_pattern=r'\b\w+\b', m
161     bigram_vect.fit(db_tr)
162     Xtr2 = bigram_vect.transform(db_tr)
163     Xte2 = bigram_vect.transform(db_te2)
164
165     print('Xtr2:(%s, %s)'%(Xtr2.shape))
166     print('Xte2:(%s, %s)%(Xte2.shape))
167
168     clf = SGDClassifier(loss='hinge', penalty='l1')
169     clf.fit(Xtr2, pola_tr)
170     pred2 = clf.predict(Xte2)
171
172     X_test2 = bigram_vect.transform(X_test)
173     print('score Unigram, test data: %s' % clf.score(Xte2, pred2))
174     print('cross validation score Unigram, test data: %s' % cross_val_score(clf,
175     print('score Bigram: %s'%clf.score(X_test2, y_test))
176     print('cross validation score Bigram: %s'%cross_val_score(clf, X_test2, y_test
177     # print(accuracy_score(y_test, pred2))
178
179     out = open('bigram.output.txt', 'w')
180     for i in range(len(pred2)):
181         out.write('%s\n'%int(pred2[i]))
182     out.close()
183
184     return Xtr2, Xte2, X_test2
185
186 ## Tf-idf
187
188 def unigram_tfidf(Xtr, Xte, X_test1, pola_tr, y_test):
189     transformer = TfidfTransformer()
190     Xtr_tfidf = transformer.fit_transform(Xtr)
191     Xte_tfidf = transformer.transform(Xte)
192
193     print('Xtr_tfidf:(%s, %s)%(Xtr_tfidf.shape))
194     print('Xte_tfidf:(%s, %s)%(Xte_tfidf.shape))
195
196     clf = SGDClassifier(loss='hinge', penalty='l1')
197     clf.fit(Xtr_tfidf, pola_tr)
198     pred3 = clf.predict(Xte_tfidf)
199
200     X_test3 = transformer.transform(X_test1)
201     print('score Unigram, test data: %s' % clf.score(Xte_tfidf, pred3))
202     print('cross validation score Unigram, test data: %s' % cross_val_score(clf,
203     print('score Unigram Tfidf: %s'%clf.score(X_test3, y_test))
204     print('cross validation score Unigram Tfidf: %s'%cross_val_score(clf, X_test3
205     # print(accuracy_score(y_test, pred3))
206
207     out = open('unigramtfidf.output.txt', 'w')
208     for i in range(len(pred3)):
209         out.write('%s\n'%int(pred3[i]))
210     out.close()

```

```

211
212     return Xtr_tfidf, Xte_tfidf, X_test3
213
214
215 ## bigram Tf-idf
216
217 def bigram_tfidf(db_tr, db_te2, X_test, pola_tr, y_test):
218     tf_vect = TfidfVectorizer(ngram_range=(1, 2), token_pattern=r'\b\w+\b', min_df=3)
219     Xtr2_tfidf = tf_vect.fit_transform(db_tr)
220     Xte2_tfidf = tf_vect.transform(db_te2)
221
222     print('Xtr2_tfidf: (%s, %s)' % (Xtr2_tfidf.shape))
223     print('Xte2_tfidf: (%s, %s)' % (Xte2_tfidf.shape))
224
225     clf = SGDClassifier(loss='hinge', penalty='l1')
226     clf.fit(Xtr2_tfidf, pola_tr)
227     pred4 = clf.predict(Xte2_tfidf)
228
229     X_test4 = tf_vect.transform(X_test)
230     print('score Unigram, test data: %s' % clf.score(Xte2_tfidf, pred4))
231     print('cross validation score Unigram, test data: %s' % cross_val_score(clf, Xtr2_tfidf, pola_tr, cv=5))
232     print('score Bigram Tfidf: %s' % clf.score(X_test4, y_test))
233     print('cross validation score Bigram Tfidf: %s' % cross_val_score(clf, X_test4, y_test, cv=5))
234     # print(accuracy_score(Xte2_tfidf, pred4))
235
236     out = open('bigramtfidf.output.txt', 'w')
237     for i in range(len(pred4)):
238         out.write('%s\n' % int(pred4[i]))
239     out.close()
240
241     return Xtr2_tfidf, Xte2_tfidf, X_test4
242
243 if __name__ == "__main__":
244     # """train a SGD classifier using unigram representation,
245     # predict sentiments on imdb_te.csv, and write output to
246     # unigram.output.txt
247     # train a SGD classifier using bigram representation,
248     # predict sentiments on imdb_te.csv, and write output to
249     # bigram.output.txt
250     # train a SGD classifier using unigram representation
251     # with tf-idf, predict sentiments on imdb_te.csv, and write
252     # output to unigramtfidf.output.txt
253     # train a SGD classifier using bigram representation
254     # with tf-idf, predict sentiments on imdb_te.csv, and write
255     # output to bigramtfidf.output.txt"""
256
257     imdb_data_preprocess(train_path)
258
259     train_data = pd.read_csv("imdb_tr.csv")
260
261     db_tr, db_te2, X_test, pola_tr, y_test = prepare_data(train_data)
262
263     Xtr, Xte, X_test1 = unigram(db_tr, db_te2, X_test, pola_tr, y_test)

```


264

265 Xtr2, Xte2, X_test2 = bigram(db_tr, db_te2, X_test, pola_tr, y_test)

266

267 Xtr_tfidf, Xte_tfidf, X_test3 = unigram_tfidf(Xtr, Xte, X_test1, pola_tr, y_test)

268

269 Xtr2_tfidf, Xte2_tfidf, X_test4 = bigram_tfidf(db_tr, db_te2, X_test, pola_tr, y_test)

270

271 end = time.clock()

272 diff = end - begin

273 **print**('time: %ss' % diff)